

One Command for All Server Environment Settings

Allen



Content

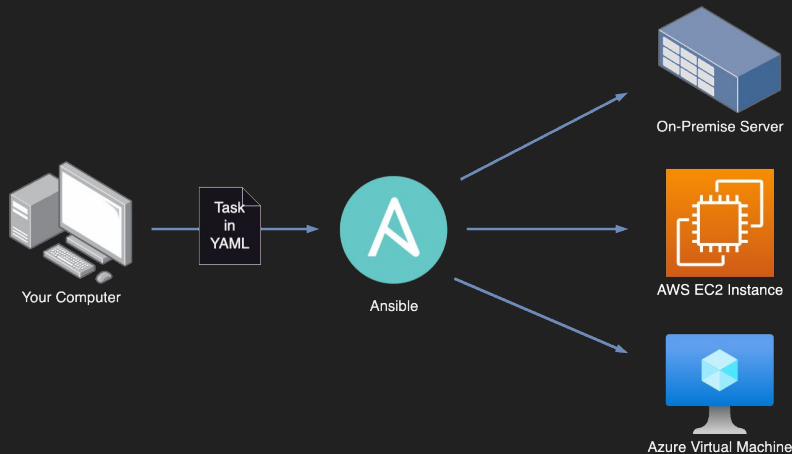
- What is Ansible?
- Inventory
- Modules
- Playbooks
- Idempotence
- Variables
- Conditionals
- Roles
- Demo

What is Ansible?

Ansible is a well-known IaC (Infrastructure as Code) tool

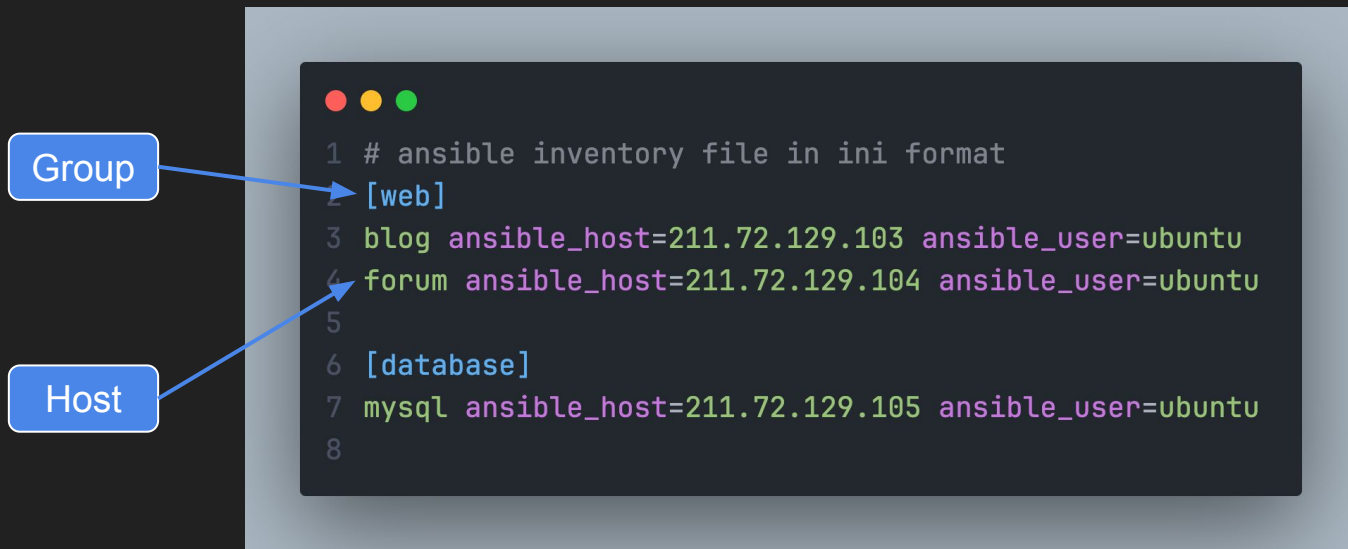
Unlike Terraform, which is used to build resources, Ansible is primarily used to configure the environment on the machine

Ansible is **Agentless**, you don't need to install any agent on the remote host to operate the remote host



Inventory

Inventory is a file used to define host connection information that Ansible refers to and connects to the remote host using SSH



P.S. Use snake_case for Group and Host names. ex. hello_world, web_server

Inventory in YAML

Besides INI, Inventory can also be written in YAML format

```
1 # ansible inventory file in yaml format
2 web:
3   hosts:
4     blog:
5       ansible_host: 211.72.129.103
6       ansible_user: ubuntu
7     forum:
8       ansible_host: 211.72.129.104
9       ansible_user: ubuntu
10 database:
11   hosts:
12     mysql:
13       ansible_host: 211.72.129.105
14       ansible_user: ubuntu
15
```

Using SSH Config with Inventory

Inventory can be used with `~/.ssh/config`

```
1 Host blog
2   User ubuntu
3   Hostname 211.72.129.103
4   Port 22
5   IdentityFile ~/.ssh/id_ed25519.pub
6
7 Host forum
8   User ubuntu
9   Hostname 211.72.129.103
10  Port 22
11  IdentityFile ~/.ssh/id_ed25519.pub
12
13 Host mysql
14   User ubuntu
15   Hostname 211.72.129.105
16   Port 22
17   IdentityFile ~/.ssh/id_ed25519.pub
18
```

```
1 # with ssh config file
2 web:
3   hosts:
4     → blog:
5     → forum:
6   database:
7     hosts:
8     → mysql:
9
```

Use Ansible to Do a Connection Test

Set up the Inventory and let's test it with the commands

- `all`: Specify the name of the target to connect to, either Group or Host. here `all` stands for all host
- `-i`: Specify the Inventory file to use
- `-m`: Specify the Ansible Module to use



```
1 ansible all -i inventory.yaml -m ping
```

Module

Ansible has a number of different modules that can be used, such as the ping module that we just used

And many other modules ...

- [apt](#): Managing apt packages on the host (for Ubuntu or Debian)
- [command](#): Execute commands on the host
- [file](#): Managing the contents of files on the host
- [fetch](#): Pulling files from the host to the local machine

Playbooks

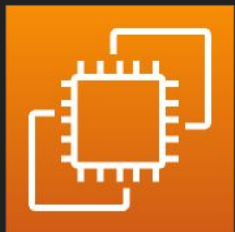
Playbooks allows you to run tasks on multiple hosts

You can even run **sequential tasks** on multiple hosts (A for B, B for C)

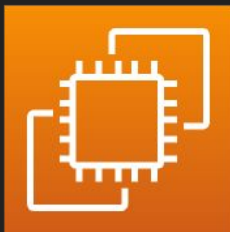
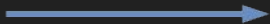
- Install Mysql
- Update the security settings
- Allow 3306 TCP Port in Firewall

- Install Laravel App
- Update Laravel settings
- Create the Database Schema

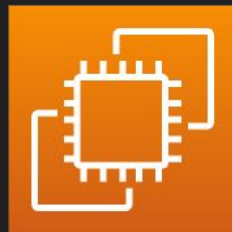
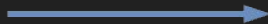
- Install Nginx
- Update the Nginx Config
- Allow 443 TCP Port in Firewall



Host A



Host B



Host C

Playbooks Example

Playbooks in YAML format

hosts: Host or Group as defined in Inventory

name: Task descriptions, Ansible Lint will require capitalization at the beginning.

module: For modules, Ansible Lint requires that they be written using fqcn (fully-qualified collection names).

parameter: Parameters of the module

```
1 - name: Install nginx and enable it
2   hosts: webserver
3   become: true
4   tasks:
5     - name: Install nginx
6       ansible.builtin.apt:
7         name: nginx
8         state: present
9     - name: Start the nginx service and enable it
10      ansible.builtin.service:
11        name: nginx
12        state: started
13        enabled: true
```

Execute Playbooks Task

Use `ansible-playbook` to run your written YAML file.



```
1 ansible-playbook playbooks.yaml -i inventory.yaml
```

Don't Specify the Inventory Every Time.

You can create an `ansible.cfg` file. This is the configuration file for Ansible, and [there are a number of settings](#) that you can set up

if you set the location of the Inventory file, you don't need to enter the command `-i inventory.yaml` afterward



```
1 [defaults]
2 inventory = ./inventory.yaml
```

Idempotence

Executing the same action one or more times will result in the same state for the object being executed

Idempotence is an important concept when writing tasks.

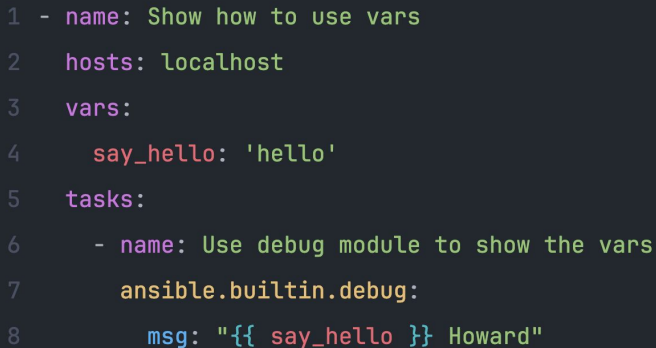
```
1 # it's ok to execute it multiple times
2 # this task will make sure the file "./foo.txt" is absent on the machine
3 - name: Install nginx and enable it
4   hosts: localhost
5   tasks:
6     - name: Remove file (delete file)
7       ansible.builtin.file:
8         path: ./foo.txt
9         state: absent
```

```
1 # each execute will append a line to the file
2 # so it's not an idempotent task
3 - name: Append a line to a file
4   hosts: localhost
5   tasks:
6     - name: Add a line to a file
7       ansible.builtin.shell:
8         cmd: echo "hello world" >> ./foo.txt
9       changed_when: true
```

Variables

You can declare a variable in playbooks so that subsequent tasks use it.

After defining the variable with `vars`, you can use the Jinja2 Filter `"{{ }}"` to convert the variable to the value you set.

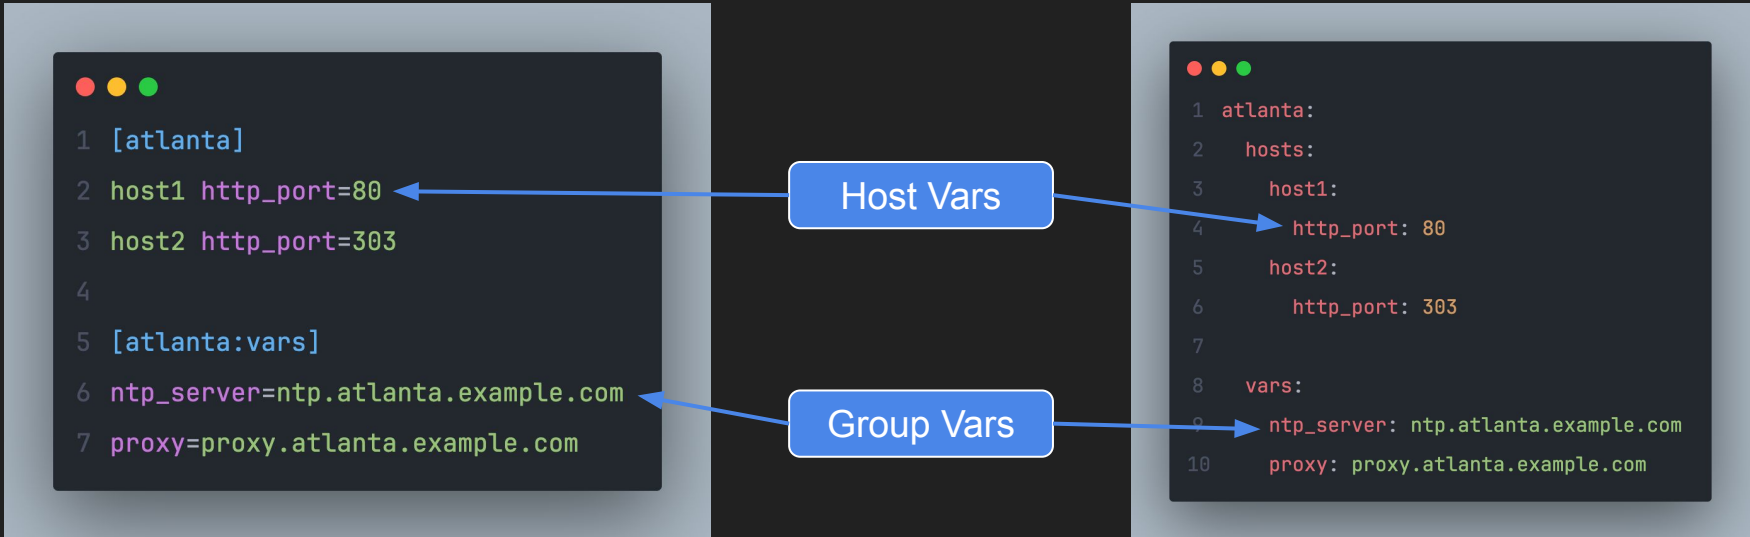
A terminal window with a dark background and light-colored text. It displays an Ansible playbook snippet. The snippet starts with a list item (line 1) defining a task named 'Show how to use vars'. It then specifies the hosts as 'localhost' (line 2). A 'vars' section (line 3) defines a variable 'say_hello' with the value 'hello' (line 4). The 'tasks' section (line 5) contains a task (line 6) that uses the 'debug' module to display the value of 'say_hello'. The message string in the debug module (line 8) uses the Jinja2 filter '{{ }}' to insert the value of 'say_hello', followed by the name 'Howard'.

```
1 - name: Show how to use vars
2   hosts: localhost
3   vars:
4     say_hello: 'hello'
5   tasks:
6     - name: Use debug module to show the vars
7       ansible.builtin.debug:
8         msg: "{{ say_hello }}" Howard"
```

You Can Set Up Variables in Inventory.

Variables can be set in the Inventory, which is divided into Group Vars and Host Vars.

Group Vars will act on all Hosts under the genus.



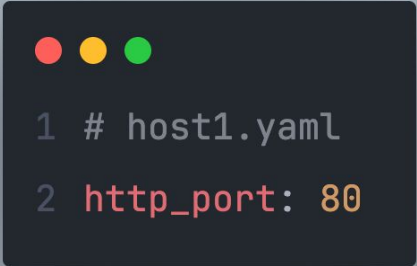
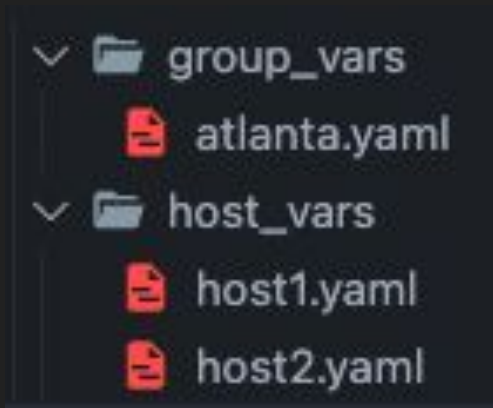
How to Use Variables in Inventory

Accessing Group Vars and Host Vars with Ansible's `hostvars`

```
1 - name: Show how to use vars in inventory
2   hosts: host1
3   tasks:
4     - name: Print the host vars 'http_port'
5       ansible.builtin.debug:
6         msg: "{{ hostvars['host1']['http_port'] }}"
7
8     - name: Print the group vars 'proxy'
9       ansible.builtin.debug:
10        msg: "{{ hostvars['host1']['proxy'] }}"
```


Organizing Your Variables with Folders

You can add a new `group_vars` and `host_vars` folder, and create a YAML file with the same name as Host or Group underneath it, and then write the variables in it.



```
1 # host1.yaml
2 http_port: 80
```

A screenshot of a terminal window with a light blue background. It shows two lines of text: a comment line `# host1.yaml` and a variable assignment `http_port: 80`. The text is color-coded: the line numbers are grey, the comment hash is grey, the variable name is red, and the value is yellow.

Register Variables

You can register the output of a task as a new variable using `register`

```
1 - name: Show how to use register
2   hosts: localhost
3   tasks:
4     - name: Register a output of a command
5       ansible.builtin.command:
6         cmd: "echo 'Hello Tavia!'"
7       changed_when: false
8       register: result
9
10    - name: Print the registered variable
11      ansible.builtin.debug:
12        msg: "{{ result.stdout }}"
13
```

Conditional

You can use the control flow in playbooks

when loop

when

When the given conditional equation holds, the task is executed.

```
1 - name: Show how to use when
2   hosts: proxy
3   tasks:
4     - name: Install Nginx on Debian
5       ansible.builtin.apt:
6         name: nginx
7         state: present
8       when: ansible_os_family == "Debian"
9
10    - name: Install Nginx on RedHat
11      ansible.builtin.yum:
12        name: nginx
13        state: present
14      when: ansible_os_family == "RedHat"
```

Loop

Bring in the values in the array and repeat the task

```
1 - name: Show how to use loop
2   hosts: localhost
3   tasks:
4     - name: Create users
5       loop:
6         - Johnson
7         - Howard
8         - Elizabeth
9         - Peggy
10        - Tavia
11      ansible.builtin.user:
12        name: "{{ item }}"
13        state: present
```

```
1 - name: Show how to use loop
2   hosts: localhost
3   tasks:
4     - name: Create users
5       loop:
6         - name: Johnson
7           uid: 1001
8         - name: Howard
9           uid: 1002
10        - name: Elizabeth
11          uid: 1003
12      ansible.builtin.user:
13        name: "{{ item.name }}"
14        uid: "{{ item.uid }}"
15        state: present
```

Role

You can modularize your tasks for reuse, similar to Terraform's Modules.

Assuming you have a playbooks installation of MySQL, you can package it as a Role, and when you later need to install MySQL on another machine, you can use this Role directly without writing it again



```
1 # create a new role 'say_hello'  
2 ansible-galaxy init say_hello
```

Role's File Structure

You can place the Role under the roles folder in the same directory as the Playbooks file

You can also place the Role in the `/etc/ansible/roles` folder



```
1 ---
2 # roles/say_hello/tasks/main.yaml
3 # tasks file for say_hello
4 - name: Say hello
5   ansible.builtin.debug:
6     msg: "Hello!"
7
```

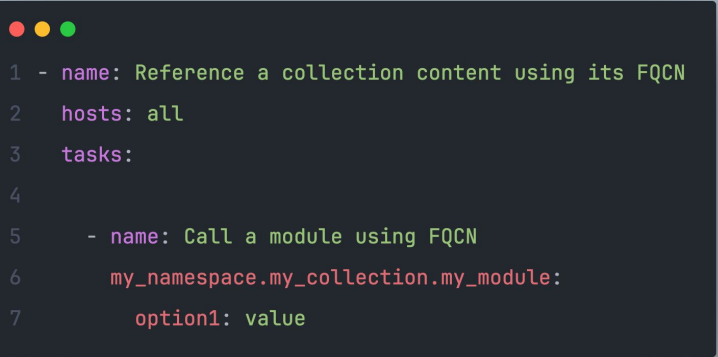
```
1 # playbooks.yaml
2 - name: Show how to use role
3   hosts: localhost
4   roles:
5     - say_hello
```

Collections

It's also a way to reuse code

Collections contain playbooks, roles, modules, and plugins

Unlike roles, which are tasks, collections describe a more complete IT infrastructure

A terminal window with a dark background and three colored window control buttons (red, yellow, green) in the top-left corner. It displays an Ansible playbook snippet with line numbers 1 through 7. The syntax is color-coded: 'name' is green, 'hosts' and 'tasks' are purple, and module names and options are red.

```
1 - name: Reference a collection content using its FQCN
2   hosts: all
3   tasks:
4
5     - name: Call a module using FQCN
6       my_namespace.my_collection.my_module:
7         option1: value
```


Ansible Galaxy

Before you start writing a new project, you can look on Ansible Galaxy to see if there is a role or collection you need.

You can also share your Role on the web for others to use.

The screenshot displays the Ansible Galaxy web interface. At the top, there's a navigation bar with the 'GALAXY' logo and links for 'About', 'Help', 'Documentation', and 'Login'. A search bar is prominently featured, showing the query 'mysql' with '(1673 results)'. Below the search bar, a 'Collections' section lists 65 results. The first result is the 'mysql' collection by 'community', which has a 4.8/5 score, 10185832 downloads, and is version 3.7.2. It lists 7 modules (mysql_db, mysql_info, mysql_query), 0 roles, and 2 plugins (mysql, mysql). The second result is 'wp_mysql_on_top_of_k8s_cluster' by 'mohit_jangir', with 25 downloads and version 1.0.0. A 'Popular Tags' sidebar on the right lists various categories with their respective counts: system (8,036), development (3,555), web (2,988), monitoring (1,803), networking (1,440), database (1,328), docker (1,317), cloud (1,210), security (1,102), and ubuntu (1,085).

Tag	Count
system	8,036
development	3,555
web	2,988
monitoring	1,803
networking	1,440
database	1,328
docker	1,317
cloud	1,210
security	1,102
ubuntu	1,085

Follow Ansible's Best Practice

[The official documentation](#) lists a number of recommendations for using Ansible and is highly recommended.

Best Practices

Here are some tips for making the most of Ansible and Ansible playbooks.

You can find some example playbooks illustrating these best practices in our [ansible-examples repository](#).
(NOTE: These may not use all of the features in the latest release, but are still an excellent reference!).

After We Talk About Ansible

I just want to say ...

Long Live the IMMUTABLE !!!